

dog_app

May 14, 2020

1 Convolutional Neural Networks

1.1 Project: Write an Algorithm for a Dog Identification App

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

Note: Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets:

Note: if you are using the Udacity workspace, you DO NOT need to re-download these - they can be found in the /data folder as noted in the cell below.

- Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location /dog_images.
- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location /lfw.

Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays human_files and dog_files.

```
In [1]: import numpy as np
        from glob import glob

        # load filenames for human and dog images
        human_files = np.array(glob("/data/lfw/*/"))
        dog_files = np.array(glob("/data/dog_images/*/"))

        # print number of images in each dataset
        print('There are %d total human images.' % len(human_files))
        print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.

There are 8351 total dog images.

Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the haarcascades directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [3]: import cv2
        import matplotlib.pyplot as plt
        %matplotlib inline

        # extract pre-trained face detector
        face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

        # load color (BGR) image
        img = cv2.imread(human_files[0])
        # convert BGR image to grayscale
        gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

        # find faces in image
        faces = face_cascade.detectMultiScale(gray)

        # print number of faces detected in the image
        print('Number of faces detected:', len(faces))
```

```

# get bounding box for each detected face
for (x,y,w,h) in faces:
    # add bounding box to color image
    cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()

```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

1.1.1 Write a Human Face Detector

We can use this procedure to write a function that returns True if a human face is detected in an image and False otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [4]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray)
    return len(faces) > 0
```

1.1.2 (IMPLEMENTATION) Assess the Human Face Detector

Question 1: Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

Answer: (You can print out your results and/or write your percentages in this cell)

```
In [5]: from tqdm import tqdm

human_files_short = human_files[:100]
dog_files_short = dog_files[:100]

#-#-# Do NOT modify the code above this line. #-#-#

## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.
count1=0
count2=0
for image in human_files_short:
    if face_detector(image):
        count1+=1
for image in dog_files_short:
    if face_detector(image):
        count2+=1
print(f"{count1}% human in human_files")
print(f"{count2}% human in dog_files")
```

```
98% human in human_files
17% human in dog_files
```

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make

use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [5]: ### (Optional)  
        ### TODO: Test performance of another face detection algorithm.  
        ### Feel free to use as many code cells as needed.
```

Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of [1000 categories](#).

```
In [6]: import torch  
        import torchvision.models as models  
  
        # define VGG16 model  
        VGG16 = models.vgg16(pretrained=True)  
  
        # check if CUDA is available  
        use_cuda = torch.cuda.is_available()  
  
        # move model to GPU if CUDA is available  
        if use_cuda:  
            VGG16 = VGG16.cuda()
```

```
Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to /root/.torch/models/vgg16-397923af.pth  
100%|| 553433881/553433881 [00:08<00:00, 68362548.10it/s]
```

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as `'dogImages/train/001.Affenpinscher/Affenpinscher_00001.jpg'`) as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```

In [7]: from PIL import Image
import torchvision.transforms as transforms

def VGG16_predict(img_path):
    """
    Use pre-trained VGG-16 model to obtain index corresponding to
    predicted ImageNet class for image at specified path

    Args:
        img_path: path to an image

    Returns:
        Index corresponding to VGG-16 model's prediction
    """
    tran=transforms.Compose([transforms.CenterCrop(244),
                             transforms.ToTensor(),
                             transforms.Normalize([0.485, 0.456, 0.406],
                                                  [0.229, 0.224, 0.225])])

    ## TODO: Complete the function.
    ## Load and pre-process an image from the given img_path
    ## Return the *index* of the predicted class for that image
    image=Image.open(img_path)
    image=tran(image)
    image=image.unsqueeze(0)
    image=image.cuda()
    output=VGG16(image).data.argmax()
    return output # predicted dex

```

1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the dog_detector function below, which returns True if a dog is detected in an image (and False if not).

```

In [8]: ### returns "True" if a dog is detected in the image stored at img_path
def dog_detector(img_path):
    ## TODO: Complete the function.
    index=VGG16_predict(img_path)
    return 150<index<269

```

1.1.6 (IMPLEMENTATION) Assess the Dog Detector

Question 2: Use the code cell below to test the performance of your dog_detector function.
 - What percentage of the images in human_files_short have a detected dog?

- What percentage of the images in `dog_files_short` have a detected dog?

Answer:

```
In [9]: ### TODO: Test the performance of the dog_detector function  
### on the images in human_files_short and dog_files_short.  
human_files_short = human_files[:100]  
dog_files_short = dog_files[:100]  
count1=0  
count2=0  
for image in human_files_short:  
    if dog_detector(image):  
        count1+=1  
for image in dog_files_short:  
    if dog_detector(image):  
        count2+=1  
print(f"{count1}% dogs in human_files")  
print(f"{count2}% dogs in dog_files")
```

0% dogs in human_files

91% dogs in dog_files

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [10]: ### (Optional)  
### TODO: Report the performance of another pre-trained network.  
### Feel free to use as many code cells as needed.
```

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel
----------	------------------------

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

Yellow Labrador	Chocolate Labrador
-----------------	--------------------

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

1.1.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dog_images/train`, `dog_images/valid`, and `dog_images/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
In [10]: import os
         from torchvision import datasets
         from PIL import ImageFile
         ImageFile.LOAD_TRUNCATED_IMAGES = True

         ### TODO: Write data loaders for training, validation, and test sets
         ## Specify appropriate transforms, and batch_sizes

         batch_size=30
         num_worker=0

         transform={
             "train": transforms.Compose([transforms.RandomRotation(45),
                                         transforms.RandomHorizontalFlip(),
                                         transforms.Resize(244),
                                         transforms.CenterCrop(244),
                                         transforms.ToTensor(),
                                         transforms.Normalize((0.485, 0.456, 0.406),
                                                             (0.229, 0.224, 0.225))]),
             "valid": transforms.Compose([transforms.Resize(244),
                                         transforms.CenterCrop(244),
                                         transforms.ToTensor(),
```



```

        transforms.Normalize((0.485, 0.456, 0.406),
                              (0.229, 0.224, 0.225))]),

    "test": transforms.Compose([transforms.Resize(244),
                                transforms.CenterCrop(244),
                                transforms.ToTensor(),
                                transforms.Normalize((0.485, 0.456, 0.406),
                                                      (0.229, 0.224, 0.225))])

}

train_data=datasets.ImageFolder("/data/dog_images/train",transform=transform["train"])
valid_data=datasets.ImageFolder("/data/dog_images/valid",transform=transform["valid"])
test_data=datasets.ImageFolder("/data/dog_images/test",transform=transform["test"])

data={
    "train":torch.utils.data.DataLoader(train_data,batch_size=batch_size,num_workers=num_
    "valid":torch.utils.data.DataLoader(valid_data,batch_size=batch_size,num_workers=num_
    "test":torch.utils.data.DataLoader(test_data,batch_size=batch_size,num_workers=num_
}

```

Question 3: Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

Answer:

I have resized images Resize(244) and CenterCrop(244). Size of input tensor is (3,224,244),
For image augmentation, i have used RandomRotation(45), RandomHorizontalFlip() and Normalize

1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```

In [54]: import torch.nn as nn
import torch.nn.functional as F

# define the CNN architecture
class Net(nn.Module):
    ### TODO: choose an architecture, and complete the class
    def __init__(self):
        super(Net, self).__init__()
        ## Define layers of a CNN
        self.conv1=nn.Conv2d(3, 32, 3, padding=1)

        self.conv2=nn.Conv2d(32, 64, 3, padding=1)
        self.bn_2=nn.BatchNorm2d(64)

        self.conv3=nn.Conv2d(64, 128, 3, padding=1)
        self.bn_3=nn.BatchNorm2d(128)

```

```

self.conv4=nn.Conv2d(128, 128, 3, padding=1)
self.bn_4=nn.BatchNorm2d(128)

self.pool=nn.MaxPool2d(2,2)
self.fc1=nn.Linear(15*15*128,512)
self.fc2=nn.Linear(512,256)
self.fc3=nn.Linear(256,133)
self.drop=nn.Dropout(0.5)

def forward(self, x):
    ## Define forward behavior
    x=self.pool(F.relu(self.conv1(x)))
    x=self.pool(F.relu(self.bn_2(self.conv2(x))))
    x=self.pool(F.relu(self.bn_3(self.conv3(x))))
    x=self.pool(F.relu(self.bn_4(self.conv4(x))))
    x=x.view(-1,15*15*128)
    x=self.drop(x)
    x=F.relu(self.fc1(x))
    x=self.drop(x)
    x=F.relu(self.fc2(x))
    x=self.drop(x)
    x=self.fc3(x)
    return x

### You so NOT have to modify the code below this line. ###

# instantiate the CNN
model_scratch = Net()

# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch.cuda()

```

Question 4: Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

Answer:

First, I like to mention that Designing Architecture for this kind of complicated task majorly depend in data argumentation.

Early version

<h4>CNN Architecture</h4>

```

<li>3 -> 8 -> 16 -> 32 -> 64 -> 128</li>
<li>3 -> 8 -> 16 -> 32 -> 64 -> 128 -> 256</li>
<li>3 -> 16 -> 32 -> 64 -> 128</li>
<li>3 -> 32 -> 64 -> 128</li>
<li>3 -> 32 -> 64 -> 128 ->256</li>
<li>3 -> 32 -> 64 -> 128 ->256 ->512</li>

```

```

    <li>3 -> 32 -> 64 -> 128 -> 256 ->256</li>
<h4>Activation function</h4>
    <li>Relu activation function is used in all layer except output layer</li>
<h4>Regularization</h4>
    <li>DropOut(0.2), DropOut(0.3), DropOut(0.4), DropOut(0.5)</li>
    <li>MaxPool2d(2,2)</li>
    <li>BatchNorm2d(depth)</li>
<h4>Fully Connected Layer</h4>
    <li>input_dim -> 500 -> 133</li>
    <li>input_dim -> 256 -> 133</li>
    <li>input_dim -> 500 -> 300 -> 133</li>
    <li>input_dim -> 512 -> 133</li>
    <li>input_dim -> 500 -> 256 -> 133</li>
<h5>Reason:</h5><p>I have try all these combination of Architecture, most of them didn't imp

```

Final version

```

<h4>CNN Architecture</h4>
    <li>3 -> 32 -> 64 -> 128 -> 128</li>
<h4>Activation function</h4>
    <li>Relu activation function is used in all layer except output layer</li>
<h4>Regularization</h4>
    <li>DropOut(0.5)</li>
    <li>MaxPool2d(2,2)</li>
    <li>BatchNorm2d(depth)</li>
<h4>Fully Connected Layer</h4>
    <li>input_dim -> 500 -> 256 -> 133</li>
<h5>Reason:</h5><p>After so many hit and trial on pervious Architecture I got maximum accuracy o

```

1.1.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```

In [55]: import torch.optim as optim

        ### TODO: select loss function
        criterion_scratch = nn.CrossEntropyLoss()

        ### TODO: select optimizer
        optimizer_scratch = optim.Adam(model_scratch.parameters(),lr=0.001)

```

1.1.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_scratch.pt'`.

```

In [56]: import time
def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
    """returns trained model"""
    # initialize tracker for minimum validation loss
    valid_loss_min = np.Inf

    for epoch in range(1, n_epochs+1):
        # initialize variables to monitor training and validation loss
        train_loss = 0.0
        valid_loss = 0.0
        start=time.time()

        #####
        # train the model #
        #####
        model.train()
        for batch_idx, (data, target) in enumerate(loaders['train']):
            # move to GPU
            if use_cuda:
                data, target = data.cuda(), target.cuda()
            ## find the loss and update the model parameters accordingly
            ## record the average training loss, using something like
            ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))
            optimizer.zero_grad()
            out=model(data)
            loss=criterion(out,target)
            loss.backward()
            optimizer.step()
            train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.item() - train_loss))
        #####
        # validate the model #
        #####
        model.eval()
        for batch_idx, (data, target) in enumerate(loaders['valid']):
            # move to GPU
            if use_cuda:
                data, target = data.cuda(), target.cuda()
            ## update the average validation loss
            out=model(data)
            loss=criterion(out,target)
            valid_loss = valid_loss + ((1 / (batch_idx + 1)) * (loss.item() - valid_loss))
        # print training/validation statistics
        print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f} \tTime Took: {}'
              .format(
                epoch,
                train_loss,
                valid_loss,
                time.time()-start
            ))

```

```

    ## TODO: save the model if validation loss has decreased
    if valid_loss<valid_loss_min:
        torch.save(model.state_dict(), save_path)
        valid_loss_min=valid_loss
        print("Saved!")
    # return trained model
    return model

# train the model
model_scratch = train(100, data, model_scratch, optimizer_scratch,
                      criterion_scratch, use_cuda, 'model_scratch.pt')

```

Epoch: 1	Training Loss: 4.941414	Validation Loss: 4.852635	Time Took: 14
Saved!			
Epoch: 2	Training Loss: 4.833686	Validation Loss: 4.785933	Time Took: 14
Saved!			
Epoch: 3	Training Loss: 4.785044	Validation Loss: 4.787337	Time Took: 14
Epoch: 4	Training Loss: 4.740670	Validation Loss: 4.695531	Time Took: 14
Saved!			
Epoch: 5	Training Loss: 4.690743	Validation Loss: 4.648328	Time Took: 14
Saved!			
Epoch: 6	Training Loss: 4.676137	Validation Loss: 4.646778	Time Took: 14
Saved!			
Epoch: 7	Training Loss: 4.647908	Validation Loss: 4.609945	Time Took: 14
Saved!			
Epoch: 8	Training Loss: 4.622295	Validation Loss: 4.597981	Time Took: 14
Saved!			
Epoch: 9	Training Loss: 4.607126	Validation Loss: 4.564602	Time Took: 14
Saved!			
Epoch: 10	Training Loss: 4.568204	Validation Loss: 4.529598	Time Took: 14
Saved!			
Epoch: 11	Training Loss: 4.549955	Validation Loss: 4.593581	Time Took: 14
Epoch: 12	Training Loss: 4.526855	Validation Loss: 4.542934	Time Took: 14
Epoch: 13	Training Loss: 4.504242	Validation Loss: 4.447408	Time Took: 14
Saved!			
Epoch: 14	Training Loss: 4.465261	Validation Loss: 4.433761	Time Took: 14
Saved!			
Epoch: 15	Training Loss: 4.451074	Validation Loss: 4.421773	Time Took: 14
Saved!			
Epoch: 16	Training Loss: 4.411558	Validation Loss: 4.366109	Time Took: 14
Saved!			
Epoch: 17	Training Loss: 4.380267	Validation Loss: 4.397091	Time Took: 14
Epoch: 18	Training Loss: 4.347218	Validation Loss: 4.543199	Time Took: 14
Epoch: 19	Training Loss: 4.281910	Validation Loss: 4.241552	Time Took: 14
Saved!			
Epoch: 20	Training Loss: 4.259721	Validation Loss: 4.197991	Time Took: 14

Saved!			
Epoch: 21	Training Loss: 4.200379	Validation Loss: 4.205253	Time Took: 1
Epoch: 22	Training Loss: 4.172917	Validation Loss: 4.158079	Time Took: 1
Saved!			
Epoch: 23	Training Loss: 4.120597	Validation Loss: 4.044804	Time Took: 1
Saved!			
Epoch: 24	Training Loss: 4.080536	Validation Loss: 4.013608	Time Took: 1
Saved!			
Epoch: 25	Training Loss: 4.049131	Validation Loss: 4.005262	Time Took: 1
Saved!			
Epoch: 26	Training Loss: 4.016993	Validation Loss: 3.932519	Time Took: 1
Saved!			
Epoch: 27	Training Loss: 3.989282	Validation Loss: 3.882300	Time Took: 1
Saved!			
Epoch: 28	Training Loss: 3.940849	Validation Loss: 3.801624	Time Took: 1
Saved!			
Epoch: 29	Training Loss: 3.906963	Validation Loss: 3.880912	Time Took: 1
Epoch: 30	Training Loss: 3.896759	Validation Loss: 3.814850	Time Took: 1
Epoch: 31	Training Loss: 3.850633	Validation Loss: 3.762826	Time Took: 1
Saved!			
Epoch: 32	Training Loss: 3.840272	Validation Loss: 3.709170	Time Took: 1
Saved!			
Epoch: 33	Training Loss: 3.776281	Validation Loss: 3.724579	Time Took: 1
Epoch: 34	Training Loss: 3.775828	Validation Loss: 3.678523	Time Took: 1
Saved!			
Epoch: 35	Training Loss: 3.728153	Validation Loss: 3.662915	Time Took: 1
Saved!			
Epoch: 36	Training Loss: 3.700731	Validation Loss: 3.627305	Time Took: 1
Saved!			
Epoch: 37	Training Loss: 3.658333	Validation Loss: 3.606213	Time Took: 1
Saved!			
Epoch: 38	Training Loss: 3.651324	Validation Loss: 3.588685	Time Took: 1
Saved!			
Epoch: 39	Training Loss: 3.644745	Validation Loss: 3.641346	Time Took: 1
Epoch: 40	Training Loss: 3.613393	Validation Loss: 3.571905	Time Took: 1
Saved!			
Epoch: 41	Training Loss: 3.577824	Validation Loss: 3.534857	Time Took: 1
Saved!			
Epoch: 42	Training Loss: 3.551511	Validation Loss: 3.576112	Time Took: 1
Epoch: 43	Training Loss: 3.531660	Validation Loss: 3.435370	Time Took: 1
Saved!			
Epoch: 44	Training Loss: 3.519716	Validation Loss: 3.445125	Time Took: 1
Epoch: 45	Training Loss: 3.479277	Validation Loss: 3.468793	Time Took: 1
Epoch: 46	Training Loss: 3.479031	Validation Loss: 3.410513	Time Took: 1
Saved!			
Epoch: 47	Training Loss: 3.454519	Validation Loss: 3.381493	Time Took: 1
Saved!			
Epoch: 48	Training Loss: 3.433129	Validation Loss: 3.399312	Time Took: 1

Epoch: 49	Training Loss: 3.394466	Validation Loss: 3.328353	Time Took: 1
Saved!			
Epoch: 50	Training Loss: 3.381346	Validation Loss: 3.344064	Time Took: 1
Epoch: 51	Training Loss: 3.360281	Validation Loss: 3.300232	Time Took: 1
Saved!			
Epoch: 52	Training Loss: 3.329239	Validation Loss: 3.245192	Time Took: 1
Saved!			
Epoch: 53	Training Loss: 3.300626	Validation Loss: 3.325883	Time Took: 1
Epoch: 54	Training Loss: 3.297437	Validation Loss: 3.252809	Time Took: 1
Epoch: 55	Training Loss: 3.281258	Validation Loss: 3.247808	Time Took: 1
Epoch: 56	Training Loss: 3.240496	Validation Loss: 3.213053	Time Took: 1
Saved!			
Epoch: 57	Training Loss: 3.215573	Validation Loss: 3.331187	Time Took: 1
Epoch: 58	Training Loss: 3.209703	Validation Loss: 3.189463	Time Took: 1
Saved!			
Epoch: 59	Training Loss: 3.221619	Validation Loss: 3.168249	Time Took: 1
Saved!			
Epoch: 60	Training Loss: 3.203685	Validation Loss: 3.211174	Time Took: 1
Epoch: 61	Training Loss: 3.181180	Validation Loss: 3.168700	Time Took: 1
Epoch: 62	Training Loss: 3.158722	Validation Loss: 3.096801	Time Took: 1
Saved!			
Epoch: 63	Training Loss: 3.135061	Validation Loss: 3.153738	Time Took: 1
Epoch: 64	Training Loss: 3.105800	Validation Loss: 3.128957	Time Took: 1
Epoch: 65	Training Loss: 3.110203	Validation Loss: 3.058874	Time Took: 1
Saved!			
Epoch: 66	Training Loss: 3.098205	Validation Loss: 3.106298	Time Took: 1
Epoch: 67	Training Loss: 3.055284	Validation Loss: 3.047528	Time Took: 1
Saved!			
Epoch: 68	Training Loss: 3.044253	Validation Loss: 3.055920	Time Took: 1
Epoch: 69	Training Loss: 3.040889	Validation Loss: 3.012601	Time Took: 1
Saved!			
Epoch: 70	Training Loss: 3.026214	Validation Loss: 3.043821	Time Took: 1
Epoch: 71	Training Loss: 3.030630	Validation Loss: 2.979015	Time Took: 1
Saved!			
Epoch: 72	Training Loss: 3.015910	Validation Loss: 3.039057	Time Took: 1
Epoch: 73	Training Loss: 2.948396	Validation Loss: 3.057590	Time Took: 1
Epoch: 74	Training Loss: 2.953296	Validation Loss: 2.991109	Time Took: 1
Epoch: 75	Training Loss: 2.933637	Validation Loss: 3.056332	Time Took: 1
Epoch: 76	Training Loss: 2.948279	Validation Loss: 3.034649	Time Took: 1
Epoch: 77	Training Loss: 2.910864	Validation Loss: 2.979034	Time Took: 1
Epoch: 78	Training Loss: 2.925085	Validation Loss: 2.940932	Time Took: 1
Saved!			
Epoch: 79	Training Loss: 2.904493	Validation Loss: 2.930520	Time Took: 1
Saved!			
Epoch: 80	Training Loss: 2.883464	Validation Loss: 2.893702	Time Took: 1
Saved!			
Epoch: 81	Training Loss: 2.853759	Validation Loss: 2.984715	Time Took: 1
Epoch: 82	Training Loss: 2.871481	Validation Loss: 2.928260	Time Took: 1

Epoch: 83	Training Loss: 2.867667	Validation Loss: 2.936177	Time Took: 1
Epoch: 84	Training Loss: 2.875197	Validation Loss: 2.908007	Time Took: 1
Epoch: 85	Training Loss: 2.848344	Validation Loss: 2.867977	Time Took: 1
Saved!			
Epoch: 86	Training Loss: 2.841463	Validation Loss: 2.865034	Time Took: 1
Saved!			
Epoch: 87	Training Loss: 2.807728	Validation Loss: 2.882562	Time Took: 1
Epoch: 88	Training Loss: 2.784172	Validation Loss: 2.947966	Time Took: 1
Epoch: 89	Training Loss: 2.764668	Validation Loss: 2.920759	Time Took: 1
Epoch: 90	Training Loss: 2.754492	Validation Loss: 2.836397	Time Took: 1
Saved!			
Epoch: 91	Training Loss: 2.757014	Validation Loss: 2.913701	Time Took: 1
Epoch: 92	Training Loss: 2.780694	Validation Loss: 2.812080	Time Took: 1
Saved!			
Epoch: 93	Training Loss: 2.736387	Validation Loss: 2.825394	Time Took: 1
Epoch: 94	Training Loss: 2.704051	Validation Loss: 2.803678	Time Took: 1
Saved!			
Epoch: 95	Training Loss: 2.707713	Validation Loss: 2.719924	Time Took: 1
Saved!			
Epoch: 96	Training Loss: 2.711509	Validation Loss: 2.844351	Time Took: 1
Epoch: 97	Training Loss: 2.703934	Validation Loss: 2.854764	Time Took: 1
Epoch: 98	Training Loss: 2.691585	Validation Loss: 2.777913	Time Took: 1
Epoch: 99	Training Loss: 2.673250	Validation Loss: 2.766931	Time Took: 1
Epoch: 100	Training Loss: 2.682129	Validation Loss: 2.883792	Time Took: 1

1.1.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```
In [58]: # load the model that got the best validation accuracy
         model_scratch.load_state_dict(torch.load('model_scratch.pt'))
         def test(loaders, model, criterion, use_cuda):

             # monitor test loss and accuracy
             test_loss = 0.
             correct = 0.
             total = 0.

             model.eval()
             for batch_idx, (data, target) in enumerate(loaders['test']):
                 # move to GPU
                 if use_cuda:
                     data, target = data.cuda(), target.cuda()
                 # forward pass: compute predicted outputs by passing inputs to the model
                 output = model(data)
                 # calculate the loss
```



```

    loss = criterion(output, target)
    # update average test loss
    test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
    # convert output probabilities to predicted class
    pred = output.data.max(1, keepdim=True)[1]
    # compare predictions to true label
    correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
    total += data.size(0)

print('Test Loss: {:.6f}\n'.format(test_loss))

print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
    100. * correct / total, correct, total))

# call test function
test(data, model_scratch, criterion_scratch, use_cuda)

```

Test Loss: 2.759789

Test Accuracy: 28% (238/836)

Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

1.1.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at dogImages/train, dogImages/valid, and dogImages/test, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```
In [11]: loaders_transfer= data
```

1.1.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable model_transfer.

```
In [21]: import torchvision.models as models
import torch.nn as nn

## TODO: Specify model architecture
```

```

model_transfer=models.vgg19(pretrained=True)

for param in model_transfer.features.parameters():
    param.requires_grad = False

n_inputs = model_transfer.classifier[6].in_features
model_transfer.classifier[6]=nn.Linear(n_inputs,133)

for param in model_transfer.classifier[6].parameters():
    param.requires_grad = True

if use_cuda:
    model_transfer = model_transfer.cuda()

```

Question 5: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer:

```

<li>I have used the vgg19 model because I noticed in style transfer exercise that vgg19 models a
<li>Model architecture is the same as vgg19 expect sixth Linear layer out is changed to 133.</li>

```

1.1.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```

In [23]: from torch import optim

criterion_transfer = nn.CrossEntropyLoss()
optimizer_transfer = optim.Adam(model_transfer.classifier[6].parameters(), lr=0.001)

```

1.1.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_transfer.pt'`.

```

In [25]: import time
def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
    """returns trained model"""
    # initialize tracker for minimum validation loss
    valid_loss_min = np.Inf

    for epoch in range(1, n_epochs+1):
        # initialize variables to monitor training and validation loss
        train_loss = 0.0
        valid_loss = 0.0
        start=time.time()

        #####

```

```

# train the model #
#####
model.train()
for batch_idx, (data, target) in enumerate(loaders['train']):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
        ## find the loss and update the model parameters accordingly
        ## record the average training loss, using something like
        ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))
        optimizer.zero_grad()
        out=model(data)
        loss=criterion(out,target)
        loss.backward()
        optimizer.step()
        train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.item() - train_loss))
    #####
# validate the model #
#####
model.eval()
for batch_idx, (data, target) in enumerate(loaders['valid']):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
        ## update the average validation loss
        out=model(data)
        loss=criterion(out,target)
        valid_loss = valid_loss + ((1 / (batch_idx + 1)) * (loss.item() - valid_loss))
# print training/validation statistics
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f} \tTime Took: {}s'.format(
    epoch,
    train_loss,
    valid_loss,
    time.time()-start
))

## TODO: save the model if validation loss has decreased
if valid_loss<valid_loss_min:
    torch.save(model.state_dict(), save_path)
    valid_loss_min=valid_loss
    print("Saved!")
# return trained model
return model

```

```
In [ ]: # train the model
```

```
    n_epochs=10
```

```
    model_transfer = train(n_epochs, loaders_transfer, model_transfer, optimizer_transfer, c
```

```
Epoch: 1
```

```
Training Loss: 1.256571
```

```
Validation Loss: 0.541658
```

```
Time Took: 33
```

Saved!

Epoch: 2	Training Loss: 0.570307	Validation Loss: 0.424150	Time Took: 31
----------	-------------------------	---------------------------	---------------

Saved!

Epoch: 3	Training Loss: 0.513456	Validation Loss: 0.428617	Time Took: 31
----------	-------------------------	---------------------------	---------------

```
In [27]: # load the model that got the best validation accuracy (uncomment the line below)
         model_transfer.load_state_dict(torch.load('model_transfer.pt'))
```

1.1.16 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
In [29]: def test(loaders, model, criterion, use_cuda):

         # monitor test loss and accuracy
         test_loss = 0.
         correct = 0.
         total = 0.

         model.eval()
         for batch_idx, (data, target) in enumerate(loaders['test']):
             # move to GPU
             if use_cuda:
                 data, target = data.cuda(), target.cuda()
             # forward pass: compute predicted outputs by passing inputs to the model
             output = model(data)
             # calculate the loss
             loss = criterion(output, target)
             # update average test loss
             test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
             # convert output probabilities to predicted class
             pred = output.data.max(1, keepdim=True)[1]
             # compare predictions to true label
             correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
             total += data.size(0)

         print('Test Loss: {:.6f}\n'.format(test_loss))

         print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
             100. * correct / total, correct, total))

         test(loaders_transfer, model_transfer, criterion_transfer, use_cuda)
```

Test Loss: 0.420989

Test Accuracy: 87% (730/836)

1.1.17 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```
In [41]: ### TODO: Write a function that takes a path to an image as input
         ### and returns the dog breed that is predicted by the model.

         # list of class names by index, i.e. a name can be accessed like class_names[0]
         class_names = [item[4:].replace("_", " ") for item in train_data.classes]

         def predict_breed_transfer(img_path):
             # load the image and return the predicted breed
             tran=transforms.Compose([transforms.CenterCrop(244),
                                     transforms.ToTensor(),
                                     transforms.Normalize([0.485, 0.456, 0.406],
                                                         [0.229, 0.224, 0.225])])

             image=Image.open(img_path)
             image=tran(image)
             image=image.unsqueeze(0)
             image=image.cuda()
             output=model_transfer(image).data.argmax()
             return class_names[output]
```

Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

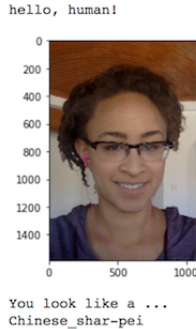
You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `human_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!

1.1.18 (IMPLEMENTATION) Write your Algorithm

```
In [56]: ### TODO: Write your algorithm.
         ### Feel free to use as many code cells as needed.

         def run_app(img_path):
             ## handle cases for a human face, dog, and neither
             dog=dog_detector(img_path)
             human = face_detector(img_path)
```



Sample Human Output

```

if(dog==True):
    print("hello, dog")
    breed = predict_breed_transfer(img_path)
    print(f"You look like a...{breed}")
    image=Image.open(img_path)
    plt.imshow(image)
    plt.show()
elif(human==True):
    print("hello, human")
    image=Image.open(img_path)
    plt.imshow(image)
    plt.show()
else:
    print("You are neither human nor dog")
    image=Image.open(img_path)
    plt.imshow(image)
    plt.show()

```

Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

1.1.19 (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

Answer: Output is better than i have expected :)

<h4>Possible points of improvement</h4>

Increasing number of epoch

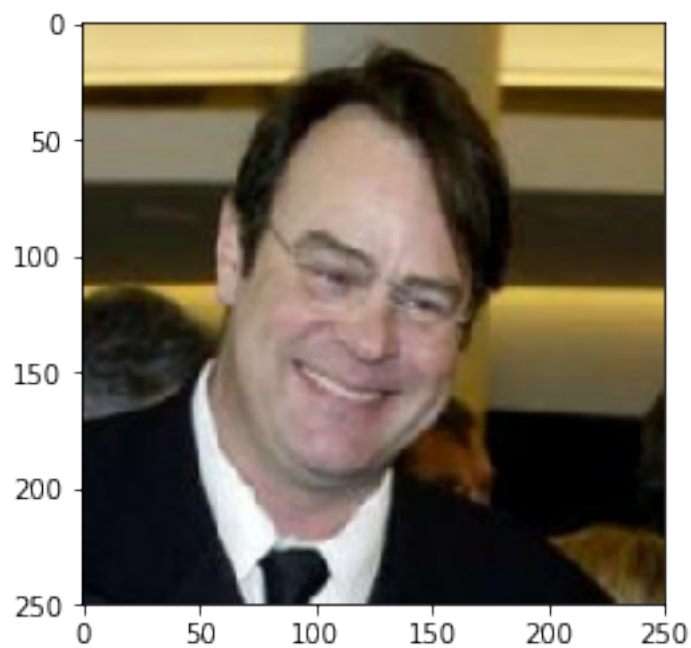
Increase number of CNN layer

Use Image augmentation to increase number of input images
DropOut2d in convolutional network can be used for regularization


```
In [57]: ## TODO: Execute your algorithm from Step 6 on
         ## at least 6 images on your computer.
         ## Feel free to use as many code cells as needed.

         ## suggested code, below
         for file in np.hstack((human_files[:3], dog_files[:3])):
             run_app(file)
```


hello, human




hello, human




hello, human




hello, dog
You look like a...German shepherd dog



hello, dog
You look like a...Mastiff



```
hello, dog  
You look like a...Mastiff
```



```
In [ ]:
```